

Party Affiliation Prediction

Michael Chen, Siyu Chen

Summary of Findings

Introduction

After some basic EDA with the stock trading data, we like to build some models that predict the party affiliation of the stockholders. This is a typical binary classification problem. The response variable is the party affiliation of the stockholders (democrat or republican). We would use Accuracy as the metric for evaluating the model because we don't value FN or FP over each other, but simply want the model that produce as many accurate predictions as possible.

Baseline Model

The model we used is a Logistic Regression model with features of transaction_year and amount. We examined in the EDA that there is correlation between party affiliation and transaction_year and amount. We used accuracy as metric for evaluating the model. For the train set, the accuracy is around 0.71. For the test set, the accuracy is 0.72.

Final Model

Two extra features are if one trade has made over 200 usd gains and the number of times ticker has been traded. They are good for data and prediction because accuracy has improved after adding these features, especially the 4th one. The final model we used was a Decision Tree with max_depth=18, min_samples_split=15, and criterion='gini' as the best parameters. The model was compared to Logistic Regression and KNN, and it outperformed them without Grid Search. Then to find the best parameters, we used Grid Search over 140 combinations of parameters.

Fairness Analysis

We want to check whether the prediction model is fair based on number of tickers when making predictions. So we manually created two group given that the number of tickers of one group is higher than average and the other group is below average. The observed difference in accuracy is 0.05678...which means that group that have higher number of tickers will also have higher accuracy when predicting. Then we did a permutation test

to see whether that phenomenon is just caused by chance. After 1000 simulations, the p-value of the permutation test is 0.035 which is lower than our significance level. So we proved that the model does not achieve accuracy parity based on number of tickers.

Code

```
In [1]: from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd
import seaborn as sns
%config InlineBackend.figure_format = 'retina' # Higher resolution figures
```

Feature Engineering

```
In [2]: merged_df = pd.read_csv('merged.csv')
merged_df = merged_df.dropna(subset=['Party'])
merged_df['Party'] = merged_df['Party'].map({'Democratic': 0, 'Republican': 1})
merged_df['transaction_year'] = merged_df['transaction_date'].apply(lambda x: x.split('-')[0])
merged_df['transaction_month'] = merged_df['transaction_date'].apply(lambda x: x.split('-')[1])
merged_df['transactions_representatives'] = merged_df.groupby('representative')['transaction_year'].transform('count')
merged_df['transactions_tickers'] = merged_df.groupby('ticker')['transaction_year'].transform('count')

merged_df
```

Out[2]:

| | disclosure_year | disclosure_date | transaction_date | owner | ticker | asset_description | type | amount | representative | district | |
|-------|-----------------|-----------------|------------------|-------|--------|---|----------|-----------------------|----------------|----------|---|
| 0 | 2021 | 2021-08-26 | 2012-06-19 | NaN | BLFSD | BioLife Solutions Inc | purchase | 1,001 — 15,000 | tom malinowski | NJ07 | https://clerk.house.gov |
| 1 | 2022 | 2022-03-03 | 2017-09-05 | NaN | SUP | Superior Industries International Inc Common S... | purchase | 1,001 — 15,000 | thomas suozzi | NY03 | https://clerk.house.gov |
| 2 | 2022 | 2022-03-03 | 2017-12-06 | NaN | CAT | Caterpillar Inc | purchase | 1,001 — 15,000 | thomas suozzi | NY03 | https://clerk.house.gov |
| 3 | 2022 | 2022-03-03 | 2018-04-17 | NaN | BA | Boeing Company | purchase | 15, 001— 50,000 | thomas suozzi | NY03 | https://clerk.house.gov |
| 4 | 2022 | 2022-03-03 | 2018-04-30 | NaN | CTRL | Control4 Corporation | purchase | 1,001 — 15,000 | thomas suozzi | NY03 | https://clerk.house.gov |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 14268 | 2022 | 2022-05-04 | 2022-04-28 | joint | MMP | Magellan Midstream Partners LP Limited Partner... | purchase | 1,001 — 15,000 | virginia foxx | NC05 | https://clerk.house.gov |
| 14269 | 2022 | 2022-05-04 | 2022-04-29 | joint | TTE | TotalEnergies Inc | purchase | 1,001 — 15,000 | virginia foxx | NC05 | https://clerk.house.gov |
| 14270 | 2022 | 2022-05-04 | 2022-04-29 | joint | ASML | ASML Holding NV - New York Registry Shares | purchase | 1,001 — 15,000 | kathy manning | NC06 | https://clerk.house.gov |

| | disclosure_year | disclosure_date | transaction_date | owner | ticker | asset_description | type | amount | representative | district | |
|-------|-----------------|-----------------|------------------|-------|--------|--|--------------|-------------------------|---------------------|----------|---|
| 14271 | 2022 | 2022-05-04 | 2022-04-29 | joint | V | Visa Inc | sale_partial | 1,001 — 15,000 | kathy manning | NC06 | https://clerk.house.gov |
| 14272 | 2022 | 2022-05-08 | 2022-05-04 | joint | NaN | Sales Tax Securitization Corp 5% Due 1/1/2027 | sale_full | 250, 001— 500,000 | suzan k. delbene | WA01 | https://clerk.house.gov |

14229 rows × 17 columns

Train Test Split

```
In [17]: x_train, x_test, y_train, y_test = train_test_split(merged_df.drop(['Party'], axis=1),\
merged_df['Party'],test_size=0.2, random_state=42)
```

Baseline Model:

Logistic Regression

Features:

- transaction_year
- amount

```
In [4]: from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline

x_train_base = x_train[['transaction_year', 'amount']]
x_test_base = x_test[['transaction_year', 'amount']]
y_train_base = y_train
y_test_base = y_test
pl = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),
               ('clf', LogisticRegression(max_iter=1000))])
pl.fit(x_train_base, y_train_base)
train_accuracy = pl.score(x_train_base, y_train_base)
y_pred_test = pl.predict(x_test_base)
test_accuracy = np.mean(y_pred_test == y_test)
print('Train accuracy:', train_accuracy)
print('Test accuracy:', test_accuracy)
```

Train accuracy: 0.7136958622507248

Test accuracy: 0.7210119465917076

Final Model

Features Tried:

- Type (no significant improvement)
- Representative, district and number of transactions by representatives cause data leakage because they expose representatives information.

Features Added:

- cap_gains_over_200_usd
- number of transactions by tickers

3rd feature: cap_gains_over_200_usd (Average Accuracy 0.01 higher than baseline)

```
In [5]: x_train_final = x_train[['transaction_year', 'amount', 'cap_gains_over_200_usd']]
x_test_final = x_test[['transaction_year', 'amount', 'cap_gains_over_200_usd']]
y_train_final = y_train
y_test_final = y_test
pl = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),
               ('clf', LogisticRegression(max_iter=1000))])
pl.fit(x_train_final, y_train_final)
train_accuracy = pl.score(x_train_final, y_train_final)
y_pred_final = pl.predict(x_test_final)
test_accuracy = np.mean(y_pred_final == y_test)
print('Train accuracy:', train_accuracy)
print('Test accuracy:', test_accuracy)
```

Train accuracy: 0.7245892998330844

Test accuracy: 0.7280393534785664

4th feature: transactions_tickers (Average Accuracy 0.03 higher than baseline)

```
In [6]: x_train_final = x_train[['transaction_year', 'amount', 'cap_gains_over_200_usd', 'transactions_tickers']]
x_test_final = x_test[['transaction_year', 'amount', 'cap_gains_over_200_usd', 'transactions_tickers']]
pl = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),
               ('clf', LogisticRegression(max_iter=1000))])
pl.fit(x_train_final, y_train_final)
train_accuracy = pl.score(x_train_final, y_train_final)
y_pred_final = pl.predict(x_test_final)
test_accuracy = np.mean(y_pred_final == y_test)
print('Train accuracy:', train_accuracy)
print('Test accuracy:', test_accuracy)
```

Train accuracy: 0.7447070192392163

Test accuracy: 0.7512297962052003

Selected Model:

Decision Tree Classifier (Average Accuracy 0.03 higher than Logistic Regression and KNN)

```
In [7]: from sklearn.tree import DecisionTreeClassifier
        from sklearn.neighbors import KNeighborsClassifier

pl = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),
               ('tree', DecisionTreeClassifier())])
pl.fit(x_train_final, y_train_final)
train_accuracy = pl.score(x_train_final, y_train_final)
y_pred_final = pl.predict(x_test_final)
test_accuracy = np.mean(y_pred_final == y_test)
print('Decision Tree Classifier:')
print('Train accuracy:', train_accuracy)
print('Test accuracy:', test_accuracy)

pl = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),
               ('tree', (KNeighborsClassifier()))])
pl.fit(x_train_final, y_train_final)
train_accuracy = pl.score(x_train_final, y_train_final)
y_pred_final = pl.predict(x_test_final)
test_accuracy = np.mean(y_pred_final == y_test)
print('\n')
print('KNeighborsClassifier:')
print('Train accuracy:', train_accuracy)
print('Test accuracy:', test_accuracy)
```

Decision Tree Classifier:
Train accuracy: 0.7852938592638145
Test accuracy: 0.7716092761770906

KNeighborsClassifier:
Train accuracy: 0.7550733550030747
Test accuracy: 0.7519325368938862

Best Parameters (Grid Search):

```
In [8]: from sklearn.model_selection import GridSearchCV
```

```
hyperparameters = {  
    'tree__max_depth': [2, 3, 4, 5, 7, 10, 13, 15, 18, None],  
    'tree__min_samples_split': [2, 3, 5, 7, 10, 15, 20],  
    'tree__criterion': ['gini', 'entropy']  
}
```

```
pl = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),  
               ('tree', DecisionTreeClassifier())])  
pl.fit(x_train_final, y_train_final)  
train_accuracy = pl.score(x_train_final, y_train_final)  
y_pred_final = pl.predict(x_test_final)  
test_accuracy = np.mean(y_pred_final == y_test)
```

```
searcher = GridSearchCV(pl, hyperparameters, cv=5)  
searcher.fit(x_train_final, y_train_final)  
print('Best parameters:', searcher.best_params_)
```

Best parameters: {'tree__criterion': 'gini', 'tree__max_depth': 18, 'tree__min_samples_split': 7}

```
In [9]: pl = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),  
                       ('tree', DecisionTreeClassifier(max_depth=18, min_samples_split=15, criterion='gini'))])  
pl.fit(x_train_final, y_train_final)  
train_accuracy_final = pl.score(x_train_final, y_train_final)  
y_pred_final = pl.predict(x_test_final)  
test_accuracy_final = np.mean(y_pred_final == y_test)  
print('Train accuracy:', train_accuracy_final)  
print('Test accuracy:', test_accuracy_final)
```

Train accuracy: 0.7737854695598699

Test accuracy: 0.7698524244553759

Fairness Analysis


```
In [11]: from sklearn import metrics
mean_tickers = merged_df.groupby('representative').count()['transactions_tickers'].mean()

results = x_test_final.copy()
results['ticker_above_avg'] = results['transactions_tickers'] > mean_tickers
results['actual'] = y_test
results['pred'] = y_pred_final
```

```
In [12]: acc_df = results.groupby('ticker_above_avg').apply(lambda x: metrics.accuracy_score(x['actual'], x['pred']))
acc_df
```

```
Out[12]: ticker_above_avg
False    0.765512
True     0.821918
dtype: float64
```

```
In [13]: obs = acc_df.diff().iloc[-1]
obs
```

```
Out[13]: 0.05640581735507455
```

Permutation Test

- Null hypothesis: The classifier's accuracy is the same for both number of tickers that is above and number of tickers that is below average, and any differences are due to chance.
- Hypothesis: The accuracy is higher for data that number of tickers is higher than average.
- Test Statistic: The difference in accuracy
- Significance Level: 0.05

```
In [19]: np.random.seed(123)
diff_in_acc = []


copy_for_shuffle = results.copy()
for _ in range(1000):

    shuffled_ticker_num = np.random.permutation(copy_for_shuffle['ticker_above_avg'].values)

    copy_for_shuffle['shuffled'] = shuffled_ticker_num

    acc = copy_for_shuffle.groupby('shuffled') \
        .apply(lambda x: metrics.accuracy_score(x['actual'], x['pred'])) \
        .diff() \
        .iloc[-1]
    diff_in_acc.append(acc)
```

```
In [20]: plt.figure(figsize=(10, 5))
pd.Series(diff_in_acc).plot(kind='hist', density=True, bins=15, title='Difference in Accuracy')
plt.axvline(x=obs, color='red', label='observed difference in accuracy')
plt.legend(loc='upper left');
```



```
In [21]: p_value = np.mean(obs <= diff_in_acc)
p_value
```

Out[21]: 0.035

Result:

Under significance level of 0.05, the p-value of the permutation test is 0.035, which will reject the null hypothesis. The result shows that there is existing unfairness in the prediction model that it does not achieve accuracy parity.

In []:

